

In sample

Labelled training split: 34,607 tweets

Teacher fine-tuned on these same tweets

Soft labels $\approx$ gold labels:
the teacher is almost always sure

Student learns only what the labels say

**No gain over fine-tuning,
whatever the committee or weighting**

Out of sample

Labelled split + up to 168,000 unlabelled
tweets the teacher has never seen

The same frozen teacher labels the new tweets

Soft labels carry the teacher's uncertainty,
which is new information

Student learns the boundary the labels
draw worst

**+0.012 macro-F1 on BERT-mini,
growing with the number of tweets**