

In sample

Labelled training split:
34,607 tweets



Teacher fine-tuned on
these same tweets



Soft labels $\approx$ gold labels:
the teacher is almost always sure



Student learns only
what the labels say



No gain over fine-tuning,
whatever the committee or weighting

Out of sample

Labelled split + 168,000 tweets
the teacher has never seen



The same frozen teacher
labels the new tweets



Soft labels carry the teacher's
uncertainty, which is new



Student learns the boundary
the labels draw worst



**+0.012 macro-F1 on BERT-mini,
growing with the number of tweets**